

Self-Supervised Learning for Video Geometry

Vanja Kovinic

1 Introduction

Recent advances in 3D vision have produced feed-forward models that reconstruct geometry without optimization. VGGT [1] predicts camera poses, depth, and structure from multi-view images in under one second; Geo4D [2] extends to dynamic scenes using video diffusion priors. Both rely on labeled or synthetic supervision, limiting scalability to billions of hours of unlabeled video online.

Videos offer an extraordinarily rich yet underutilized signal for geometric learning. Unlike static images, videos provide: (1) natural multi-view geometry through camera motion, (2) temporal consistency constraints that implicitly encode 3D structure, (3) motion cues that reveal depth ordering and object boundaries, and (4) virtually unlimited scale - billions of hours of video exist without geometric labels. Self-supervised learning transformed language [3] and vision [4, 5] by exploiting unlabeled data. Can it do the same for video geometry?

Self-supervised monocular depth exists [6], but methods focus on per-frame predictions. Designing objectives that teach how 3D structure evolves over time remains unexplored. Which temporal self-supervised objectives best capture 4D structure? How data-efficient is video pretraining? What geometric understanding emerges from unlabeled videos?

2 Background

2.1 Feed-Forward Multi-View Reconstruction

Multi-view reconstruction traditionally requires optimization-based structure-from-motion [7]. Recent work replaces this with feed-forward prediction. DUST3R [8] pioneered dense point regression from image pairs but requires global alignment for multiple views. VGGT [1] extends this with camera tokens and alternating attention, processing 1-200+ images jointly. MonST3R [9] performs pairwise tracking for dynamic scenes; Geo4D [2] adapts video diffusion models. All require expensive supervision: DUST3R/VGGT use structure-from-motion labels, Geo4D uses synthetic data rendered from 3D assets, limiting real-world scalability.

2.2 Self-Supervised Learning

Self-supervised learning extracts training signals from data structure. Masked autoencoders [4] reconstruct masked patches; DINO [5, 10] matches predictions between augmented views, producing semantic features without labels. VideoMAE [11] extends masking to spatiotemporal tubes. For geometry, self-supervised depth [6, 12] uses photometric consistency between consecutive frames. However, methods predict per-frame depth independently, treating video as image sequences rather than modeling temporal geometric reasoning. Whether self-supervised objectives can teach spatiotemporal structure (tracking evolving 3D geometry and maintaining multi-view consistency) remains unexplored, yet is fundamental to scaling video geometric understanding.

3 Problem Statement

While self-supervised learning has transformed image understanding, its application to video geometry remains limited by several fundamental challenges:

Appearance vs geometry - Geometric learning requires different signals than appearance-based pretraining. While methods like DINO learn semantic features through augmentation invariance, geometry demands precise metric structure and multi-view consistency. Models can minimize photometric loss through surface appearance matching without capturing underlying 3D structure. Do appearance-based objectives transfer to geometric reasoning, or is geometry-specific supervision necessary?

Temporal signals and shortcuts - Videos encode geometric constraints through motion parallax and occlusion patterns, yet naive temporal objectives permit trivial solutions - copying pixels across frames or learning optical flow without depth understanding. Which temporal self-supervised objectives force genuine spatiotemporal geometric reasoning rather than enabling such shortcuts?

Scaling to real video - Current methods use expensive supervision (VGGT) or synthetic data (Geo4D), limiting scale. Whether pretraining on real unlabeled video can match these approaches, how much data is required, and which objectives transfer most effectively to downstream tasks remain uncharacterized - yet are essential for scaling video geometric learning.

References

- [1] Jianyuan Wang et al. *VGGT: Visual Geometry Grounded Transformer*. 2025. arXiv: 2503.11651 [cs.CV]. URL: <https://arxiv.org/abs/2503.11651>.
- [2] Zeren Jiang et al. *Geo4D: Leveraging Video Generators for Geometric 4D Scene Reconstruction*. 2025. arXiv: 2504.07961 [cs.CV]. URL: <https://arxiv.org/abs/2504.07961>.
- [3] Jacob Devlin et al. *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*. 2019. arXiv: 1810.04805 [cs.CL]. URL: <https://arxiv.org/abs/1810.04805>.
- [4] Kaiming He et al. *Masked Autoencoders Are Scalable Vision Learners*. 2021. arXiv: 2111.06377 [cs.CV]. URL: <https://arxiv.org/abs/2111.06377>.
- [5] Kaiming He et al. *Masked Autoencoders Are Scalable Vision Learners*. 2021. arXiv: 2111.06377 [cs.CV]. URL: <https://arxiv.org/abs/2111.06377>.
- [6] Clément Godard et al. *Digging Into Self-Supervised Monocular Depth Estimation*. 2019. arXiv: 1806.01260 [cs.CV]. URL: <https://arxiv.org/abs/1806.01260>.
- [7] Carlo Tomasi and Takeo Kanade. “Shape and motion from image streams under orthography: a factorization method”. In: *Int. J. Comput. Vision* 9.2 (Nov. 1992), pp. 137–154. ISSN: 0920-5691. DOI: 10.1007/BF00129684. URL: <https://doi.org/10.1007/BF00129684>.
- [8] Shuzhe Wang et al. *DUST3R: Geometric 3D Vision Made Easy*. 2024. arXiv: 2312.14132 [cs.CV]. URL: <https://arxiv.org/abs/2312.14132>.
- [9] Junyi Zhang et al. *MonST3R: A Simple Approach for Estimating Geometry in the Presence of Motion*. 2025. arXiv: 2410.03825 [cs.CV]. URL: <https://arxiv.org/abs/2410.03825>.
- [10] Maxime Oquab et al. *DINOv2: Learning Robust Visual Features without Supervision*. 2024. arXiv: 2304.07193 [cs.CV]. URL: <https://arxiv.org/abs/2304.07193>.
- [11] Zhan Tong et al. *VideoMAE: Masked Autoencoders are Data-Efficient Learners for Self-Supervised Video Pre-Training*. 2022. arXiv: 2203.12602 [cs.CV]. URL: <https://arxiv.org/abs/2203.12602>.
- [12] Jamie Watson et al. *The Temporal Opportunist: Self-Supervised Multi-Frame Monocular Depth*. 2021. arXiv: 2104.14540 [cs.CV]. URL: <https://arxiv.org/abs/2104.14540>.